

Hierarchical Width-Based Planning and Learning

Miquel Junyent, Vicenç Gómez, Anders Jonsson

Universitat Pompeu Fabra, Barcelona, Spain

Background: IW(w)


- **Breadth-first search** (BrFS)
 - States factored into **features** $\phi(s)$
 - **Prunes states** that are **not novel**
 - A state is **novel** if it has a **feature tuple of size w** that is **new in the search**
 - **Complexity exponential in w** , but independent of $|S|$
 - Most classical planning benchmarks present a **low width** with **single atom goals**
- Example for $w=1$ with 3 binary features and 4 actions

# Domains	# Inst.	Inst. IW(1)	Inst. IW(2)
37	37,921	37.0%	88.2%

- In practice:


37	37,921	37.0%	88.2
----	--------	-------	------

 - Problems have **higher width** (no single-atom goal tasks)
 - IW(w) is mostly used with **w=1** due to computational constraints



Hierarchical IW (HIW)

- Blind search methods require two components:
 - **Successor function:** given a state and an action, returns a successor state (e.g., simulator)
 - **Stopping condition:** tells us when to stop the search (e.g., goal is met)
 - Our **hierarchical approach to blind search:**
 - Considers two sets of features F_h and F_ℓ
 - Modifies:
 - **High-level successor function:** each call triggers a low-level search
 - **Low-level stopping condition:** stops the low-level search when a state s that maps to a different $\phi_h(s)$ is found
 - We can pause and **resume low level searches**
 - Allows for **many levels of abstraction**
 - Accepts **different planners** at each level
- Q = Queue(root)
 While Q not empty:
 s = PopFirst(Q)
 For each action a:
 x = GenerateSuccessor(s, a)
 Append(Q, x)
 If ShouldStop(x):
 return
- Example BRFS
- IW(2)



Hierarchical IW

- **IW(w) at the different levels**
- For instance:
 - IW(2) at high-level
 - IW(1) at low-level
- In general: $\text{HIW}(w_h, w_\ell)$
- HIW can solve problems of width $w_h + w_\ell$



- Features:
- 1-D position (**low level**)
 - Having the key (**high level**)

This problem has **width 2**,
but it can be solved by
HIW(1,1)

Complexity Results

- Let $N(n, d, w)$ denote the **maximum amount of novel nodes** that IW(w) generates in a problem with $n = |F|$ features of domain size $d = |D|$
- Two basic premises:
 - A feature **has one value** at a time
 - A feature value **appears in several tuples** simultaneously

- Recursive formula:

$$N(n, d, 0) = 1, \longrightarrow \text{Only the initial state is novel}$$
$$N(n, d, n) = d^n, \longrightarrow \text{All states are novel}$$

$$N(n, d, w) = (d-1)N(n-1, d, w-1) + N(n-1, d, w).$$

States novel due to **one feature** f


States novel due to **other features**

- General formula, for $0 < w < n$:

$$N(n, d, w) = \sum_{k=0}^w \left[\binom{n-1-k}{w-k} d^k (d-1)^{w-k} \right]$$

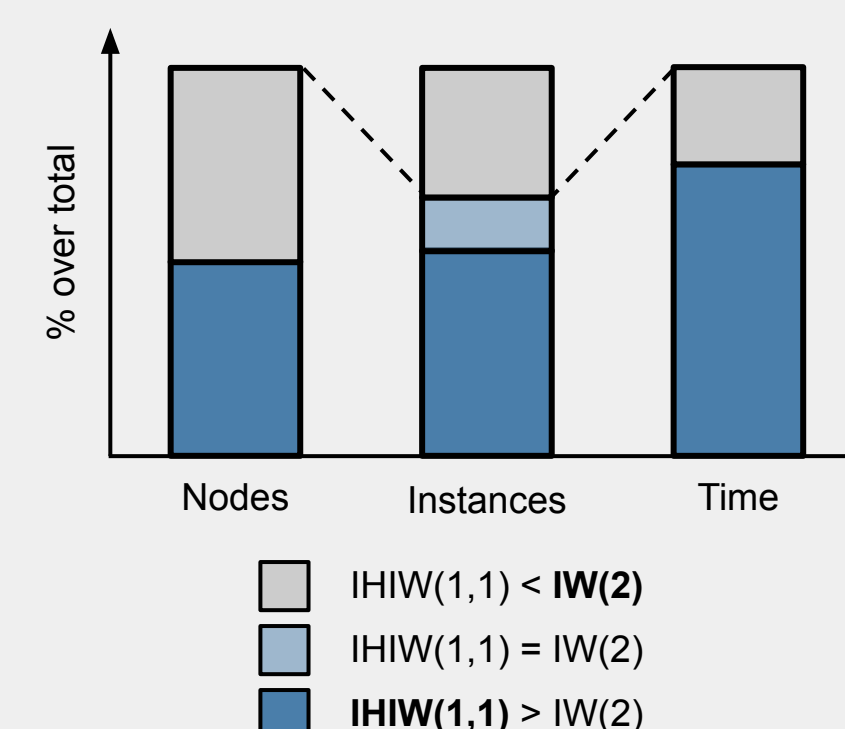
HIW in Classical Planning

- **Hypothesis:** features that **only change once** in a branch before being pruned are good candidates
- **Incremental HIW(1,1):**
 - Iteratively **run HIW(1,1)**
 - **Add one feature** to F_h at each iteration
 - **Discover new features** when necessary
 - **Reuse** the search **tree** among iterations



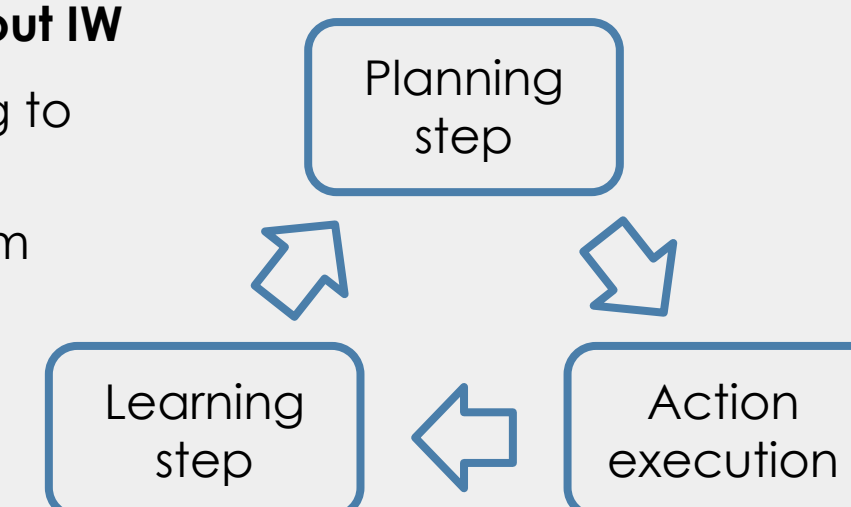
Results in Classical Planning

- Single goal instances
- Budget of 10K nodes
- We report:
 - Solved instances (%)
 - Avg. nodes (solved)
 - Avg. time in s (solved)
- **IHIW > IW(1) in 31/36 domains**
- Compared to IW(2):
 - **>= in 24/36 domains**
 - Uses **less nodes in 12/24**
 - Solves it **faster in 18/24**


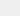
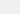


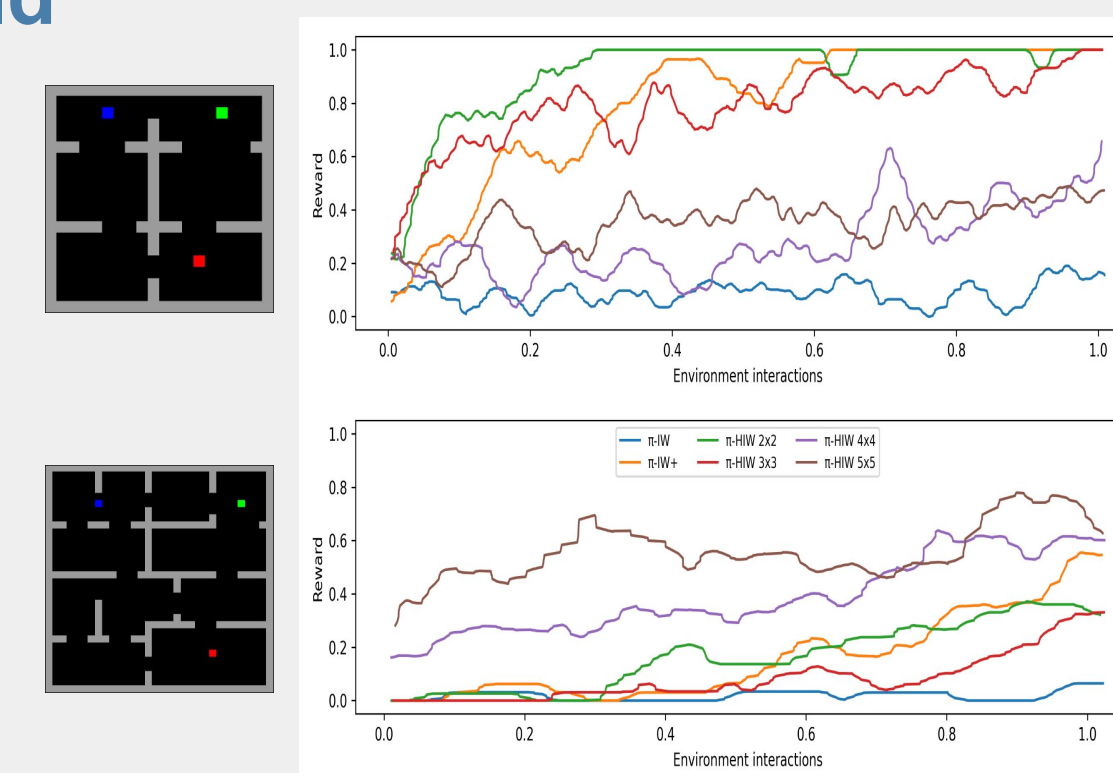
- π -HIW: Planning & Learning

- We integrate HIW with a **policy learning** scheme:
 - **High-level** planner: **Count-Based Rollout IW**
 - Selects high-level nodes according to $p \propto \exp(1/\tau(c+1))$
 - Prunes nodes using a mapping from novel tuples to unpruned nodes
 - **Low-level** planner: **π -IW modified**
 - Tree counts for tie-breaking
 - Value function





Results in gridworld


- Two sparse reward tasks
- The episode terminates:
 - when the agent  picks the key  and reaches the door  ($r = +1$)
 - when hitting a wall ($r = -1$)
 - after 200 / 500 steps ($r = 0$)



Results in Atari games



- Neural network activations **(low level)**
- Downsampling **(high level)**



miquel.junyent@upf.edu